

A Programmer's Blog

about C/C++, Java, Linux ...

Notes for playing with ptrace on 64 bits Ubuntu 12.10

 Wei Xu / January 12, 2013 / code reading, linux

This blog is the notes during I learning the “Playing with ptrace”(<http://www.linuxjournal.com/article/6100>).

The original examples was using 32 bits machine, which doesn't work on my 64 bits Ubuntu 12.10.

Let's start from the first ptrace example:

```
1  #include <sys/ptrace.h>
2  #include <sys/types.h>
3  #include <sys/wait.h>
4  #include <unistd.h>
5  #include <linux/user.h>    /* For constants
6                               ORIG_EAX etc */
7  int main()
8  {   pid_t child;
9       long orig_eax;
10      child = fork();
11      if (child == 0) {
12          ptrace(PTRACE_TRACEME, 0, NULL, NULL);
13          execl( "/bin/ls" , "ls" , NULL);
14      }
15      else {
16          wait(NULL);
17          orig_eax = ptrace(PTRACE_PEEKUSER,
18                           child, 4 * ORIG_EAX,
19                           NULL);
20          printf( "The child made a "
21                 "system call %ldn" , orig_eax);
22          ptrace(PTRACE_CONT, child, NULL, NULL);
```

```
23     }
24     return 0;
25 }
```

The compiler shows the following error:

```
1 fatal error: 'linux/user.h' file not found
2 #include <linux/user.h>
```

Something need to change because of:

1. The 'linux/user.h' no longer exists
2. The 64 bits register is R*X, so EAX changed to RAX

There are two solutions to fix this:

1. change 'linux/user.h' to 'sys/reg.h', and use:

```
long original_rax = ptrace(PTRACE_PEEKUSER, child, 8 * ORIG_RAX, NULL);
```

The addr changed from '**4 * ORIG_EAX**' to '**8 * ORIG_RAX**' because it's the address to read in the user area, and the orig_rax member in user_regs_struct is the **15th** member(start from 0). The definition of ORIG_RAX in file 'sys/reg.h' specify it's position: **# define ORIG_RAX 15**. Because of the other members has size 8 on 64 bits machine, so the addr is: 8 * ORIG_RAX.

The definition of struct user_regs_struct and user in file 'sys/user.h':

```
1 struct user_regs_struct
2 {
3     unsigned long int r15;
4     unsigned long int r14;
5     unsigned long int r13;
6     unsigned long int r12;
7     unsigned long int rbp;
8     unsigned long int rbx;
9     unsigned long int r11;
10    unsigned long int r10;
```

```
11     unsigned long int r9;
12     unsigned long int r8;
13     unsigned long int rax;
14     unsigned long int rcx;
15     unsigned long int rdx;
16     unsigned long int rsi;
17     unsigned long int rdi;
18     unsigned long int orig_rax;
19     unsigned long int rip;
20     unsigned long int cs;
21     unsigned long int eflags;
22     unsigned long int rsp;
23     unsigned long int ss;
24     unsigned long int fs_base;
25     unsigned long int gs_base;
26     unsigned long int ds;
27     unsigned long int es;
28     unsigned long int fs;
29     unsigned long int gs;
30 };
31
32 struct user
33 {
34     struct user_regs_struct    regs;
35     int        u_fpvalid;
36     struct user_fpregs_struct i387;
37     unsigned long int    u_tsize;
38     unsigned long int    u_dsize;
39     unsigned long int    u_ssize;
40     unsigned long int    start_code;
41     unsigned long int    start_stack;
42     long int    signal;
43     int        reserved;
44     struct user_regs_struct*    u_ar0;
45     struct user_fpregs_struct*    u_fpstate;
46     unsigned long int    magic;
47     char        u_comm [32];
```

```
48     unsigned long int    u_debugreg [8];
49 };
```

2. change 'linux/user.h' to 'sys/user.h', and use

```
1 struct user_regs_struct regs;
2 ptrace(PTRACE_GETREGS, child, NULL, &regs);
3 printf("The child made a system call %ld\n", regs.orig_rax);
```

The second one is simpler because it doesn't need to calculate the position, but it read more data than the first one.

I think it would be more clear and easier to understand if we use the address of the orig_rax field directly:

```
1 struct user* user_space = (struct user*)0;
2 long original_rax = ptrace(PTRACE_PEEKUSER, child, &user_space-
>regs.orig_rax, NULL);
```

We can compile and run it now, but we got: '**The child made a system call 59**', which is different with '**11**' from the original post, is there anything wrong? From the file sys/syscall.h, it included file 'asm/unistd.h' and the comment says that file list the system calls:

```
1 /* This file should list the numbers of the system calls the system
2 knows.
3 But instead of duplicating this we use the information available
4 from the kernel sources. */
#include <asm/unistd.h>
```

The file 'asm/unistd.h' include different files based on __i386__ and __ILP32__:

```
1 # ifdef __i386__
2 #  include <asm/unistd_32.h>
3 # elif defined(__ILP32__)
4 #  include <asm/unistd_x32.h>
```

```
5 # else
6 # include <asm/unistd_64.h>
7 # endif
```

From the file ‘asm/unistd_64.h’ which contains the system call names for 64 bits machine, we can found that:

```
#define __NR_execve 59
```

Ok, that’s all for the first example, and after understand it, it’s easy to understand the rest parts in part I(<http://www.linuxjournal.com/article/6100>) and part II(<http://www.linuxjournal.com/article/6210>).

Share this:



Like this:

Loading...

 Wei Xu / January 12, 2013 / code reading, linux

5 thoughts on “Notes for playing with ptrace on 64 bits Ubuntu 12.10”



Pira

March 27, 2013 at 21:41

This is useful for me , thank you!



Romain

October 24, 2013 at 20:03

Got exactly the same problem following the same article (quite obviously), this helps me! Thanks



jiaxi he

August 25, 2016 at 11:14

meet register problem in x64, this help me a lot



Durgesh

September 5, 2016 at 01:28

Nice piece of information, really helped me a lot



sagar suman

August 1, 2018 at 05:14

This article was really helpful. In 64 bit systems function calling convention has been changed(arguments passed to the registers). This should also be addressed in the article.